

Project: Investigate The Movie Database (TMDb)

February 25, 2019

1 Project: Investigate The Movie Database (TMDb)

--by Lu Tang

1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

Introduction

Dataset: This data set contains information about 10,000 movies collected from The Movie Database (TMDb), including user ratings and revenue. Data can be download from [here](#). The final two columns ending with “_adj” show the budget and revenue of the associated movie in terms of 2010 dollars, accounting for inflation over time.

The project aims to explore the following questions: - Question 1: What are the trend for movie industry? Are movie industry making more money over years - Question 2: Are newer movies more popular? - Question 3: What kinds of properties are associated with movies that have high revenues? - Question 4. Is it possible to make extremely high profit movies with low budget? - Question 5: What are the top 10 rated movies? and how is their profitability?

```
In [36]: # import library that will be used in this project
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Data Wrangling

1.1.1 General Properties

```
In [37]: # loading data
tmdb=pd.read_csv('tmdb-movies.csv')
# show number of rows and columns
print(tmdb.shape)
```

```
# to avoid truncated output
pd.options.display.max_columns = 150
# show first 2 rows
tmdb.head(2)
```

(10866, 21)

```
Out[37]:
```

	id	imdb_id	popularity	budget	revenue	original_title
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road

	cast
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...

	homepage	director	tagline
0	http://www.jurassicworld.com/	Colin Trevorrow	The park is open.
1	http://www.madmaxmovie.com/	George Miller	What a Lovely Day.

	keywords
0	monster dna tyrannosaurus rex velociraptor island
1	future chase post-apocalyptic dystopia australia

	overview	runtime
0	Twenty-two years after the events of Jurassic ...	124
1	An apocalyptic story set in the furthest reach...	120

	genres
0	Action Adventure Science Fiction Thriller
1	Action Adventure Science Fiction Thriller

	production_companies	release_date	vote_count
0	Universal Studios Amblin Entertainment Legenda...	6/9/15	5562
1	Village Roadshow Pictures Kennedy Miller Produ...	5/13/15	6185

	vote_average	release_year	budget_adj	revenue_adj
0	6.5	2015	1.379999e+08	1.392446e+09
1	7.1	2015	1.379999e+08	3.481613e+08

Initial observation: - Our focus will be analyzing movie properties associated with high revenue, some columns are irrelevant for our analysis, e.g id,imdb_id, homepage, tagline, keywords, overview, production_companies, release_date (since we already have release_year). - we can also remove budget and revenue, since we have budget_adj and revenue_adj to analyze.

```
In [38]: # Drop extraneous columns
drop_col=['id','imdb_id','homepage','tagline','keywords','overview','production_compan
tmdb = tmdb.drop(drop_col, axis=1)
```

```
# check the result
tmdb.head(1)
```

```
Out[38]:
```

	popularity	original_title	cast	director	runtime	genres	vote_count	vote_average	release_year	budget_adj	revenue_adj
0	32.985763	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	124	Action Adventure Science Fiction Thriller	5562	6.5	2015	1.379999e+08	1.392446e+09

```
In [39]: # check data type and missing values
tmdb.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 11 columns):
popularity      10866 non-null float64
original_title  10866 non-null object
cast            10790 non-null object
director        10822 non-null object
runtime         10866 non-null int64
genres          10843 non-null object
vote_count      10866 non-null int64
vote_average    10866 non-null float64
release_year    10866 non-null int64
budget_adj      10866 non-null float64
revenue_adj     10866 non-null float64
dtypes: float64(4), int64(3), object(4)
memory usage: 933.9+ KB
```

```
In [40]: # check statistical information
tmdb.describe(include='all')
```

```
Out[40]:
```

	popularity	original_title	cast	director	runtime
count	10866.000000	10866	10790	10822	10866.000000
unique	NaN	10571	10719	5067	NaN
top	NaN	Hamlet	Louis C.K.	Woody Allen	NaN
freq	NaN	4	6	45	NaN
mean	0.646441	NaN	NaN	NaN	102.070863
std	1.000185	NaN	NaN	NaN	31.381405
min	0.000065	NaN	NaN	NaN	0.000000
25%	0.207583	NaN	NaN	NaN	90.000000
50%	0.383856	NaN	NaN	NaN	99.000000

75%	0.713817	NaN	NaN	NaN	111.000000
max	32.985763	NaN	NaN	NaN	900.000000

	genres	vote_count	vote_average	release_year	budget_adj \
count	10843	10866.000000	10866.000000	10866.000000	1.086600e+04
unique	2039	NaN	NaN	NaN	NaN
top	Drama	NaN	NaN	NaN	NaN
freq	712	NaN	NaN	NaN	NaN
mean	NaN	217.389748	5.974922	2001.322658	1.755104e+07
std	NaN	575.619058	0.935142	12.812941	3.430616e+07
min	NaN	10.000000	1.500000	1960.000000	0.000000e+00
25%	NaN	17.000000	5.400000	1995.000000	0.000000e+00
50%	NaN	38.000000	6.000000	2006.000000	0.000000e+00
75%	NaN	145.750000	6.600000	2011.000000	2.085325e+07
max	NaN	9767.000000	9.200000	2015.000000	4.250000e+08

	revenue_adj
count	1.086600e+04
unique	NaN
top	NaN
freq	NaN
mean	5.136436e+07
std	1.446325e+08
min	0.000000e+00
25%	0.000000e+00
50%	0.000000e+00
75%	3.369710e+07
max	2.827124e+09

Insights: - Some columns contain NaN values, but the amount is not significant; we don't need to drop all the nulls at the beginning. - Data type are all correct. - The minimum runtime is 0, which is impossible, and some movies have extremely long runtime, we will investigate the outlier data - budget_adj and revenue_adj have minimum and median value as 0 too, which is odd, and the difference from 75% to maximum is huge, we need to investigate in the later analysis process - popularity, vote_count has very uneven distribution, with some extreme high value data.

1.1.2 Data Cleaning

This dataset is generally clean, column names are also clear and with preferred snakecase. For some string columns that contains '|', we will clean and analyze in the later part specific to the question we want to answer.

1. Remove duplicated data

```
In [41]: # check how many rows are duplicated
sum(tmdb.duplicated())
```

```
Out[41]: 1
```

```
In [42]: # Drop duplicated rows
tmdb.drop_duplicates(inplace=True)

# douch check the results
sum(tmdb.duplicated())
```

Out[42]: 0

2. Cleaning abnormal data for runtime

```
In [43]: # Find out how many rows are 0 for runtime
sum(tmdb["runtime"]==0)
```

Out[43]: 31

```
In [44]: # Since it is impossible to have runtime as 0, we will remove these.
tmdb=tmdb[tmdb["runtime"]>0]

#double check the result
sum(tmdb["runtime"]==0)
```

Out[44]: 0

3. Cleaning abnormal data for budget

```
In [45]: sum(tmdb["budget_adj"]==0)
```

Out[45]: 5668

```
In [46]: # It is impossible to make a movie without any budget, we will remove these data
tmdb=tmdb[tmdb["budget_adj"]>0]

# Double check the result
sum(tmdb["budget_adj"]==0)
```

Out[46]: 0

4. Cleaning abnormal data for revenue

```
In [47]: sum(tmdb["revenue_adj"]==0)
```

Out[47]: 1312

```
In [48]: # It is impossible to make a movie without any budget, we will remove these data
tmdb=tmdb[tmdb["revenue_adj"]>0]

# Double check the result
sum(tmdb["revenue_adj"]==0)
```

Out[48]: 0

```
In [49]: # Double check the cleaning result
print(tmdb.shape)
tmdb.head(1)
```

```
(3854, 11)
```

```
Out[49]: popularity original_title \
0 32.985763 Jurassic World

cast director \
0 Chris Pratt|Bryce Dallas Howard|Irrfan Khan|Vi... Colin Trevorrow

runtime genres vote_count \
0 124 Action|Adventure|Science Fiction|Thriller 5562

vote_average release_year budget_adj revenue_adj
0 6.5 2015 1.379999e+08 1.392446e+09
```


```
## Exploratory Data Analysis
```

1.2 1. Find pattern and visualize relationship

1.1. Explore relations with revenue_adj

```
In [50]: # plot a heatmap to see correlation with `revenue_adj` for each columns
plt.figure(figsize=(8,5))
sns.heatmap(tmdb.corr(),annot=True,cmap='coolwarm')
plt.xticks(rotation=45)
plt.title('Correlation heatmap for whole movie data')
```

```
Out[50]: Text(0.5, 1.0, 'Correlation heatmap for whole movie data')
```



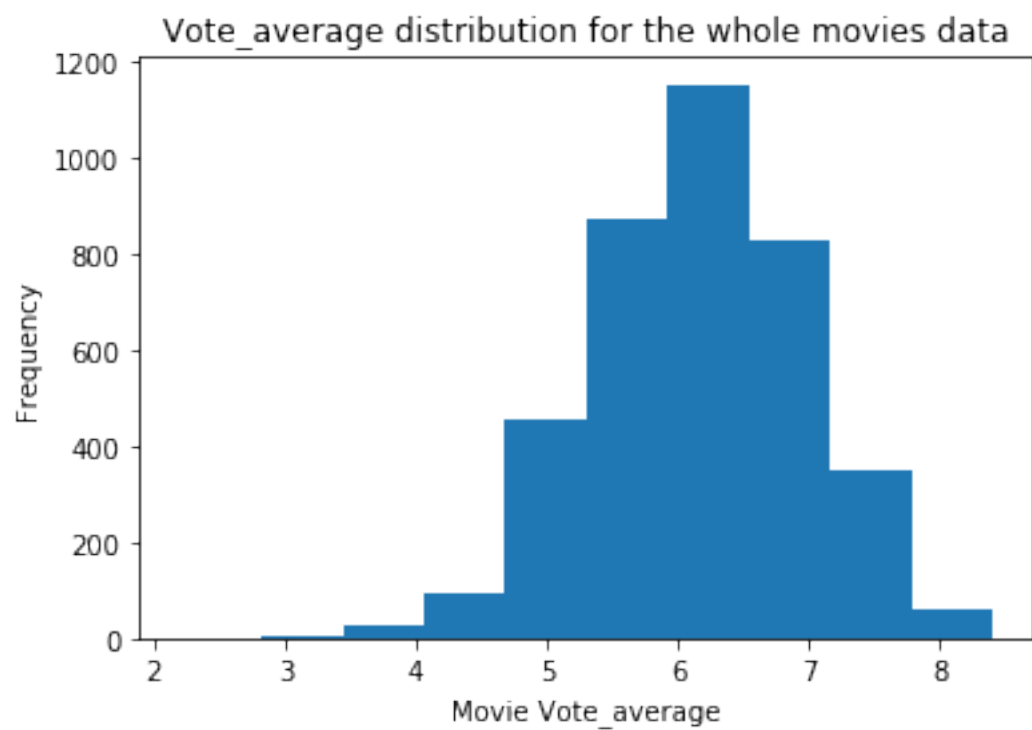
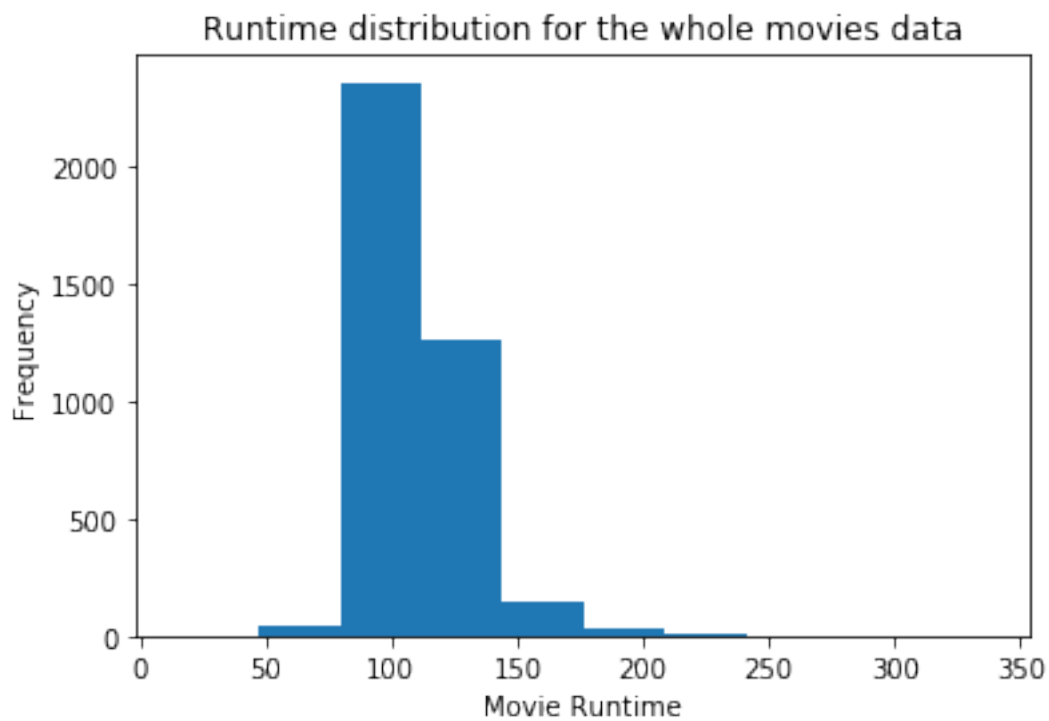
Conclusion: >- revenue_adj is positive-correlated with popularity, vote_count and budget_adj, which makes sense, the more popular, the more vote_count and and more revenues. And high budget movies are expected with high revenue too. >- popularity and vote_count are strongly correlated eith each other. >- runtime, vote_average and release_year do not have strong relation with any other columns. In fact release_year is slightly negative-correlated with revenue_adj.


1_2. Plotting charts to find out the distribution for the variables that do not have strong correlation with Movie Revenue, i.e. runtime, vote_average, release_year.

```
In [51]: # plotting distribution for 'runtime'
tmdb['runtime'].plot.hist(title='Runtime distribution for the whole movies data')
plt.xlabel('Movie Runtime')
plt.show()

# plotting distribution for 'vote_average'
tmdb['vote_average'].plot.hist(title='Vote_average distribution for the whole movies')
plt.xlabel('Movie Vote_average')
plt.show()

# plotting distribution for 'release_year'
tmdb['release_year'].plot.hist(title=('Movie Release_year Distribution'))
plt.show()
```






Conclusion: >- Most movies have median length from about 100 minutes to 180 minutes.
>- 'vote_average' has normal distribution, with average around 6. >- There are more movies produced over time.

1_3. Plotting scatter chart to explore detailed relationship between popularity and vote_count, and find out outliers.

```
In [52]: # plotting relation for 'popularity' and 'vote_count'
         tmdb.plot.scatter(x='popularity',y='vote_count')
```

```
Out[52]: <matplotlib.axes._subplots.AxesSubplot at 0x1a16b7b860>
```






Conclusion: >- From the scatter chart, we can confirm popularity and vote_count have strong positive correlation, same result as from the heatmap; however, we can also notice there are three movies rated extremely high popularity, but vote count is not extremely high. >- If we run regression model to decide movie revenues, we have to choose of one of them as an independent variable, but this is beyond the goal of this project.

1_4. Plotting scatter charts to further explore the relation with revenue_adj for the variables of popularity, vote_count and budget_adj.

```
In [53]: tmdb.plot.scatter(x='popularity', y='revenue_adj')
tmdb.plot.scatter(x='vote_count', y='revenue_adj')
tmdb.plot.scatter(x='budget_adj', y='revenue_adj')
```

```
Out[53]: <matplotlib.axes._subplots.AxesSubplot at 0x1a15516940>
```






Conclusion: >- In general, the three variables (popularity, vote_count and budget_adj) are all positively correlated with revenue_adj, but the correlation is not very strong, which is the same conclusion from the heatmap; >- There are many outlier data, some movies with extremely high popularity and high vote_count do not have extremely high revenue. These movies may be controversial, and popularity and vote_count alone are not good indicators for movie success. >- Also, some extremely high budget movies do not have very high revenue, which means they may be losing money.

1.3 2. Explore Answers for research questions

1.3.1 Question 1. What are the profitability trend for movie industry?

```
In [54]: 2. # create a column for profit
          tmdb['profit'] = tmdb['revenue_adj'] - tmdb['budget_adj']
          tmdb['profit'].plot.box()
```

```
Out[54]: <matplotlib.axes._subplots.AxesSubplot at 0x1a16a39588>
```



- Some movies are losing money; others, however have huge profit.

```
In [55]: # Create a dataframe that group by year and calculate mean value
year_mean=tmdb.groupby('release_year').mean()
year_mean.head(2)
```


```
Out [55]:
```

	popularity	runtime	vote_count	vote_average	budget_adj \
release_year					
1960	1.324513	130.0	372.6	7.40	3.068179e+07
1961	0.787718	132.5	191.4	6.62	2.818516e+07

	revenue_adj	profit
release_year		
1960	1.902299e+08	1.595481e+08
1961	2.463622e+08	2.181770e+08

```
In [56]: # plotting line chart for profit
plt.plot(year_mean.index, year_mean['profit'],label='Profit')
plt.xlabel('years')
plt.ylabel('In terms of 2010 dollars')
plt.title('Profit over years')
plt.legend()
```

```
Out [56]: <matplotlib.legend.Legend at 0x1a15605550>
```

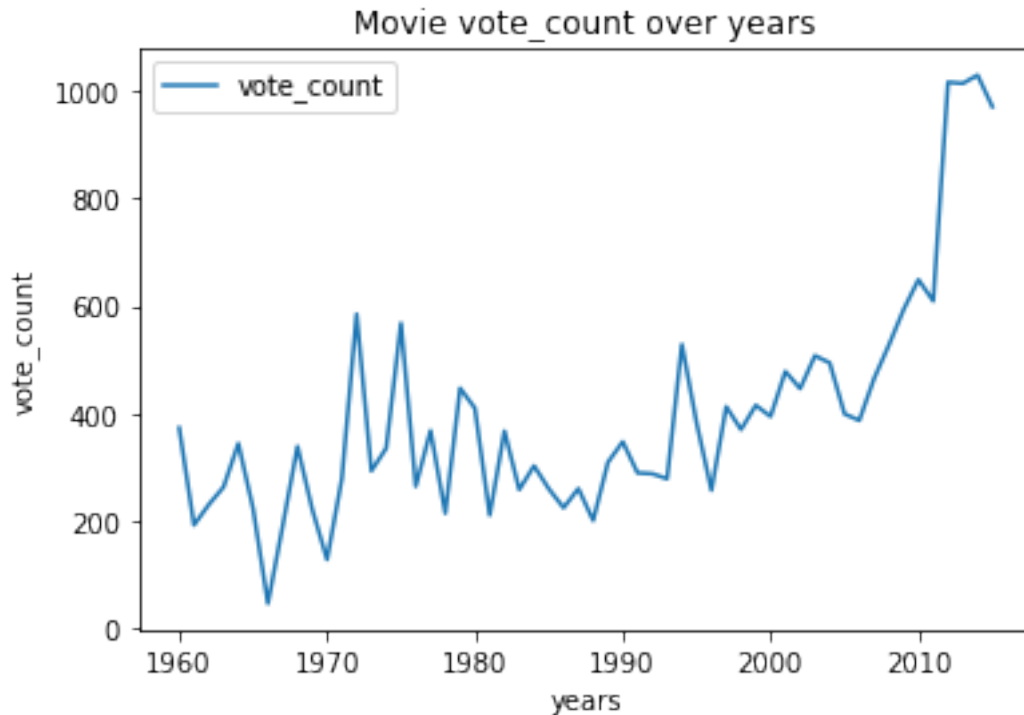


Conclusion: >- Overall, the average profit per year is lower since 1980; There is less profit to making a movie compared with three decades ago. >- In the earlier years from 1960 to 1980, film industry have higher profit but with very high fluctuation too, and the profit trend is more stable in recent years. We can conclude in the earlier years, film industry is relatively new, high risk is associated with high profit.

1.3.2 Question 2: Are newer movies more popular?

```
In [57]: # Since popularuty variable has some outlier data, we use vote_count to count for pop
plt.plot(year_mean.index, year_mean['vote_count'], label='vote_count')
plt.xlabel('years')
plt.ylabel('vote_count')
plt.title('Movie vote_count over years')
plt.legend()
```

```
Out[57]: <matplotlib.legend.Legend at 0x1a16c10da0>
```



Conclusion: >- Yes. There is clear trend that the newer movies are more popular.

1.3.3 Question 3: What kinds of properties are associated with movies that have high revenues?

1. Find out the movies with very high revenues

```
In [58]: # check the distribution for 'revenue_adj'
tmdb.revenue_adj.describe([.8,.9]).iloc[3:]
```

```
Out[58]: min      2.370705e+00
          50%      6.173068e+07
          80%      2.033811e+08
          90%      3.538761e+08
          max      2.827124e+09
          Name: revenue_adj, dtype: float64
```

```
In [59]: # Create an ordinal data column to categorize movies with different levels of revenue.
bin_edge=[0, 2.872138e+07, 1.496016e+08, 2.880722e+08, 3e+09]
bin_names=['very_low', 'low', 'high', 'very_high']
tmdb['revenue_level']=pd.cut(tmdb.revenue_adj,bin_edge,labels=bin_names)
tmdb['revenue_level'].value_counts()
```

```
Out[59]: low      1550
          very_low  1271
```

```

very_high      517
high           516
Name: revenue_level, dtype: int64

```

- Half of the movies have very_low or even negative revenue
- There are **517 movies with very high_revenue**, let's focus on these movies.

```

In [60]: # Create a dataframe that only contains movies with very high revenues
very_high=tmdb[tmdb['revenue_level']=='very_high']
# View top 5 very_high revenue movies
very_high.head()

```

```

Out[60]:
   popularity  original_title \
0    32.985763      Jurassic World
1    28.419936      Mad Max: Fury Road
3    11.173104  Star Wars: The Force Awakens
4     9.335014      Furious 7
5     9.110700      The Revenant

   cast \
0  Chris Pratt|Bryce Dallas Howard|Irrfan Khan|Vi...
1  Tom Hardy|Charlize Theron|Hugh Keays-Byrne|Nic...
3  Harrison Ford|Mark Hamill|Carrie Fisher|Adam D...
4  Vin Diesel|Paul Walker|Jason Statham|Michelle ...
5  Leonardo DiCaprio|Tom Hardy|Will Poulter|Domhn...

   director  runtime \
0    Colin Trevorrow    124
1    George Miller    120
3    J.J. Abrams    136
4    James Wan    137
5  Alejandro González Iñárritu    156

   genres  vote_count  vote_average \
0  Action|Adventure|Science Fiction|Thriller    5562      6.5
1  Action|Adventure|Science Fiction|Thriller    6185      7.1
3  Action|Adventure|Science Fiction|Fantasy    5292      7.5
4  Action|Crime|Thriller    2947      7.3
5  Western|Drama|Adventure|Thriller    3929      7.2

   release_year  budget_adj  revenue_adj  profit  revenue_level
0          2015  1.379999e+08  1.392446e+09  1.254446e+09  very_high
1          2015  1.379999e+08  3.481613e+08  2.101614e+08  very_high
3          2015  1.839999e+08  1.902723e+09  1.718723e+09  very_high
4          2015  1.747999e+08  1.385749e+09  1.210949e+09  very_high
5          2015  1.241999e+08  4.903142e+08  3.661143e+08  very_high

```

2. Find out among the very_high profit movies, what kinds of genres are the most common ones.


```
In [61]: # separate the genres with '/' and make it a new table
genres=very_high['genres'].str.split('|', expand=True)
# view the new table
genres.head()
```

```
Out[61]:
```

	0	1	2	3	4
0	Action	Adventure	Science Fiction	Thriller	None
1	Action	Adventure	Science Fiction	Thriller	None
3	Action	Adventure	Science Fiction	Fantasy	None
4	Action	Crime	Thriller	None	None
5	Western	Drama	Adventure	Thriller	None

```
In [62]: # Count the frequency for each genre for column 0, sorted as index.
col_0=genres.loc[:,0].value_counts().sort_index()
col_0
```

```
Out[62]:
```

Action	121
Adventure	116
Animation	44
Comedy	63
Crime	12
Drama	61
Family	12
Fantasy	26
History	3
Horror	9
Music	6
Mystery	3
Romance	6
Science Fiction	21
Thriller	9
War	3
Western	2

Name: 0, dtype: int64

```
In [63]: # Convert the pandas Series to a data frame
df_0=pd.DataFrame(data=col_0, index=col_0.index)
df_0
```

```
Out[63]:
```

	0
Action	121
Adventure	116
Animation	44
Comedy	63
Crime	12
Drama	61
Family	12
Fantasy	26
History	3

Horror	9
Music	6
Mystery	3
Romance	6
Science Fiction	21
Thriller	9
War	3
Western	2

```
In [64]: # do the same for other columns:
col_1=generes.loc[:,1].value_counts().sort_index()
df_1=pd.DataFrame(data=col_1, index=col_1.index)

col_2=generes.loc[:,2].value_counts().sort_index()
df_2=pd.DataFrame(data=col_2, index=col_2.index)

col_3=generes.loc[:,3].value_counts().sort_index()
df_3=pd.DataFrame(data=col_3, index=col_3.index)

col_4=generes.loc[:,4].value_counts().sort_index()
df_4=pd.DataFrame(data=col_4, index=col_4.index)

In [65]: # join the other 4 dataframe together
generes_join=df_0.join(df_1).join(df_2).join(df_3).join(df_4)
generes_join.head()
```

```
Out[65]:
```

	0	1	2	3	4
Action	121	79	33	3.0	4.0
Adventure	116	74	36	14.0	1.0
Animation	44	23	9	3.0	2.0
Comedy	63	44	37	13.0	5.0
Crime	12	16	19	11.0	2.0

```
In [66]: generes_join.fillna(0)
```

```
Out[66]:
```

	0	1	2	3	4
Action	121	79	33	3.0	4.0
Adventure	116	74	36	14.0	1.0
Animation	44	23	9	3.0	2.0
Comedy	63	44	37	13.0	5.0
Crime	12	16	19	11.0	2.0
Drama	61	52	32	5.0	1.0
Family	12	36	40	28.0	9.0
Fantasy	26	42	23	15.0	10.0
History	3	5	4	1.0	0.0
Horror	9	5	3	1.0	1.0
Music	6	6	3	2.0	4.0
Mystery	3	13	7	10.0	2.0
Romance	6	29	20	10.0	5.0

Science Fiction	21	18	37	32.0	6.0
Thriller	9	38	66	28.0	7.0
War	3	1	11	4.0	1.0
Western	2	2	3	0.0	0.0

```
In [67]: # calculate the sum of each genre's frequency
genres_join['sum_genres']=genres_join.sum(axis=1)
```

```
In [69]: # sort value based on frequency sum and show top 5
genres_join.sort_values('sum_genres', ascending=False).head()
```

```
Out[69]:
```

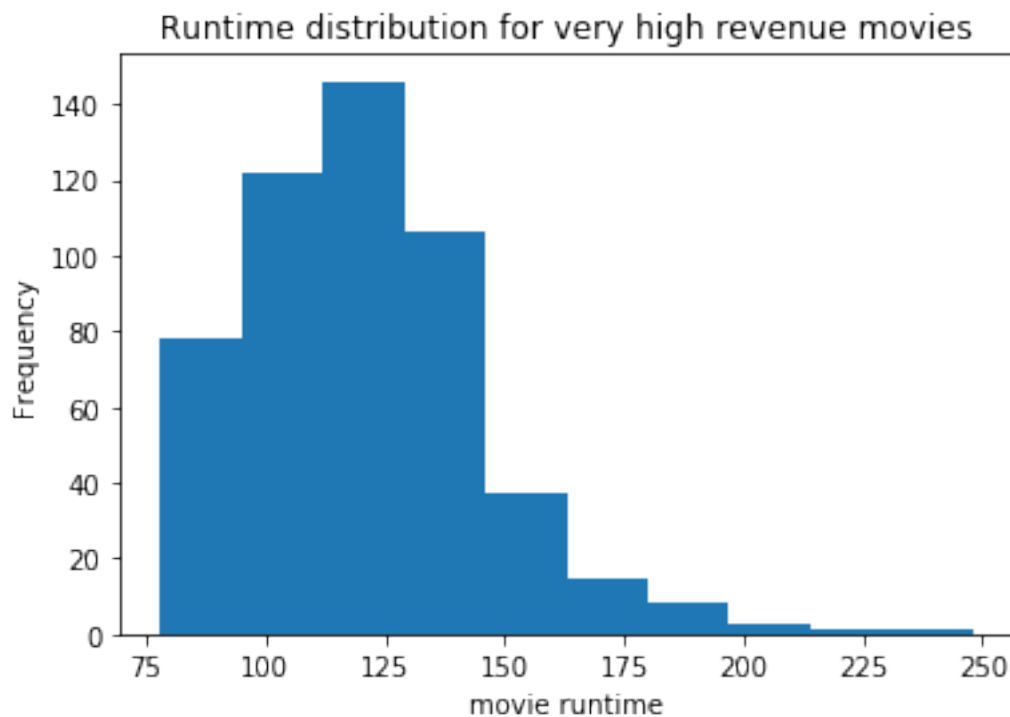
	0	1	2	3	4	sum_genres
Adventure	116	74	36	14.0	1.0	241.0
Action	121	79	33	3.0	4.0	240.0
Comedy	63	44	37	13.0	5.0	162.0
Drama	61	52	32	5.0	1.0	151.0
Thriller	9	38	66	28.0	7.0	148.0

- Now we can see among these very_high revenue movies, the genres with 'Action', 'Adventure', 'Comedy', 'Drama', 'Thriller' are top five genres

3. Explore movie runtime for very_high revenue movies

```
In [70]: # check the runtime distribution
very_high['runtime'].plot.hist()
plt.xlabel('movie runtime')
plt.title('Runtime distribution for very high revenue movies')
```

```
Out[70]: Text(0.5, 1.0, 'Runtime distribution for very high revenue movies')
```

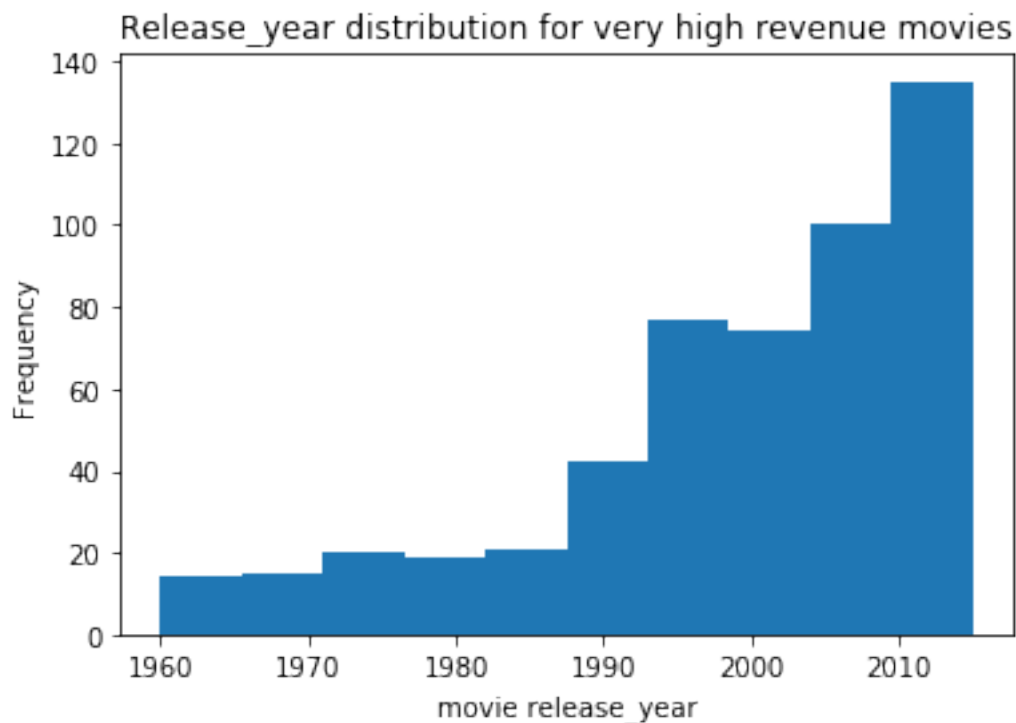


- Most movies in this dataset have runtime ranging from **75 to 150 minutes, with 100 to 125 the most popular.**
- Although some movies that have long runtime also have high revenue, a movie made under 1 hour is less likely to have very high revenue.

4. Explore release_year for very_high revenue movies

```
In [71]: # check the release_year distribution
very_high['release_year'].plot.hist()
plt.xlabel('movie release_year')
plt.title('Release_year distribution for very high revenue movies')
```

Out[71]: Text(0.5, 1.0, 'Release_year distribution for very high revenue movies')

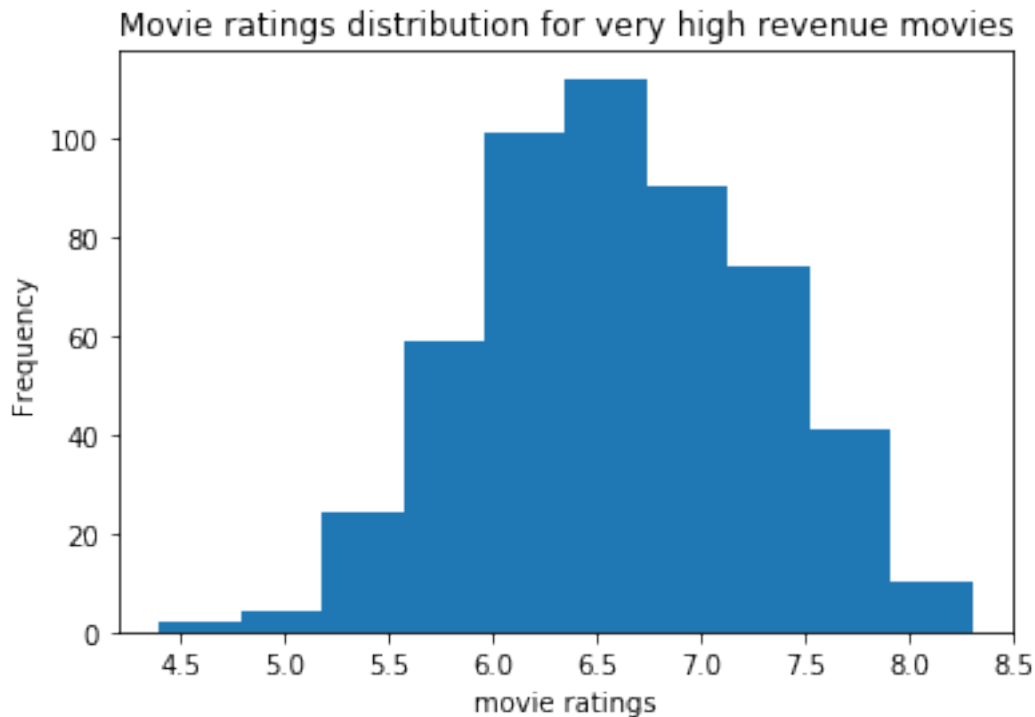


- The movies in this dataset range from 1960 to 2015, which is very similar distribution with the plot from whole dataset.
- Since there are more movies produced over years and from the line chart previously, we can conclude newer movies does not mean higher profit, it will depend on the movie itself.

5. Explore movie ratings for very_high revenue movies

```
In [72]: # check the vote_average' distribution
very_high['vote_average'].plot.hist()
plt.xlabel('movie ratings')
plt.title('Movie ratings distribution for very high revenue movies')
```

```
Out[72]: Text(0.5, 1.0, 'Movie ratings distribution for very high revenue movies')
```




- The chart reveals similar pattern compared with the plotting for whole dataset, but with higher rating in general ,ranging from 4.5 to 8.5, and average is roughly 6.5. From the heatmap earlier, we also see there isn't strong relation between vote_average and high revenue, but a movie with low rating will not have very high revenue.

6. Explore budget for very_high revenue movies

```
In [73]: # Are these high revenue movies made with high budget?
very_high.plot.scatter(x='revenue_adj', y='budget_adj')
plt.title('Relation between revenue and budget')
```

```
Out[73]: Text(0.5, 1.0, 'Relation between revenue and budget')
```



- From the heatmap for movies with whole dataset, we found budget and revenue are positively related with each other, however, as we can see among these high_revenue movies, the correlation pattern is weak.
- There are some high revenue movies with low budget and vice versa

1.3.4 Question 4. Is it possible to make extremely high profit movies with low budget?

```
In [74]: # find out the median value of budget for the whole data
tmdb.budget_adj.median()
```

```
Out[74]: 30016111.9054567
```

```
In [75]: # create a table for low_budget movies
low_budget=tmdb[tmdb['budget_adj']<30016111.9054567]
# check how many rows are in this table
print(low_budget.shape[0])
# view the distribution of 'revenue_level' in low_budget movie table
low_budget['revenue_level'].value_counts()
```

```
1927
```

```
Out[75]: very_low    1020
         low         746
```

```

high          100
very_high     61
Name: revenue_level, dtype: int64

```

```

In [76]: # Percentage of very_high revenue movies with low_budget
61/1927

```

```

Out[76]: 0.0316554229372081

```

- There are only 0.03 of the movies in the low budget group but with very_high revenue, i.e most low budget movies does not yeild high revenue, but it does not mean it is imposible.

```

In [89]: # Create a table for these movies
low_budget_very_high_revenue=low_budget[low_budget['revenue_level']=='very_high']
# check how many rows are in this table
print(low_budget_very_high_revenue.shape[0])
low_budget_very_high_revenue.head(3)

```

```

61

```

```

Out[89]:
      popularity  original_title \
1340    0.602862  Saturday Night Fever
1403    2.367474    The Blind Side
1922    5.293180    Black Swan

      cast      director \
1340  John Travolta|Karen Lynn Gorney|Barry Miller|J...  John Badham
1403  Sandra Bullock|Quinton Aaron|Kathy Bates|Tim M...  John Lee Hancock
1922  Natalie Portman|Mila Kunis|Vincent Cassel|Barb...  Darren Aronofsky

      runtime  genres  vote_count  vote_average  release_year \
1340     118  Drama|Music      192         6.3         1977
1403     129    Drama     1078         7.1         2009
1922     108  Drama|Mystery|Thriller     2597         7.1         2010

      budget_adj  revenue_adj  profit revenue_level
1340  1.259223e+07  8.530813e+08  8.404891e+08    very_high
1403  2.947561e+07  3.142795e+08  2.848038e+08    very_high
1922  1.300000e+07  3.278037e+08  3.148037e+08    very_high

```

```

In [90]: # calculate the profit mean value in the category
low_budget_very_high_revenue['profit'].mean()

```

```

Out[90]: 519202694.07224876

```

```

In [91]: # We can also sort profit from the whole data to see the top 10 most profitable movie.
top_10_profit=tmdb.sort_values('profit',ascending=False).head(10)
top_10_profit

```

```

Out[91]:
popularity      original_title \
1329    12.037933      Star Wars
1386     9.432768      Avatar
5231     4.355219      Titanic
10594    2.010733      The Exorcist
9806     2.563191      Jaws
8889     2.900556      E.T. the Extra-Terrestrial
3        11.173104      Star Wars: The Force Awakens
8094     1.136610      The Net
10110    2.631987      One Hundred and One Dalmatians
7309     5.488441      The Empire Strikes Back

cast \
1329    Mark Hamill|Harrison Ford|Carrie Fisher|Peter ...
1386    Sam Worthington|Zoe Saldana|Sigourney Weaver|S...
5231    Kate Winslet|Leonardo DiCaprio|Frances Fisher|...
10594    Linda Blair|Max von Sydow|Ellen Burstyn|Jason ...
9806    Roy Scheider|Robert Shaw|Richard Dreyfuss|Lorr...
8889    Henry Thomas|Drew Barrymore|Robert MacNaughton...
3        Harrison Ford|Mark Hamill|Carrie Fisher|Adam D...
8094    Sandra Bullock|Jeremy Northam|Dennis Miller|We...
10110    Rod Taylor|J. Pat O'Malley|Betty Lou Gerson|Ma...
7309    Mark Hamill|Harrison Ford|Carrie Fisher|Billy ...

director runtime \
1329      George Lucas      121
1386      James Cameron      162
5231      James Cameron      194
10594     William Friedkin      122
9806      Steven Spielberg      124
8889      Steven Spielberg      115
3          J.J. Abrams      136
8094      Irwin Winkler      114
10110     Clyde Geronimi|Hamilton Luske|Wolfgang Reitherman      79
7309      Irvin Kershner      124

genres  vote_count  vote_average \
1329      Adventure|Action|Science Fiction      4428      7.9
1386      Action|Adventure|Fantasy|Science Fiction      8458      7.1
5231      Drama|Romance|Thriller      4654      7.3
10594      Drama|Horror|Thriller      1113      7.2
9806      Horror|Thriller|Adventure      1415      7.3
8889      Science Fiction|Adventure|Family|Fantasy      1830      7.2
3          Action|Adventure|Science Fiction|Fantasy      5292      7.5
8094      Crime|Drama|Mystery|Thriller|Action      201      5.6
10110      Adventure|Animation|Comedy|Family      913      6.6
7309      Adventure|Action|Science Fiction      3954      8.0

```


	release_year	budget_adj	revenue_adj	profit	revenue_level
1329	1977	3.957559e+07	2.789712e+09	2.750137e+09	very_high
1386	2009	2.408869e+08	2.827124e+09	2.586237e+09	very_high
5231	1997	2.716921e+08	2.506406e+09	2.234714e+09	very_high
10594	1973	3.928928e+07	2.167325e+09	2.128036e+09	very_high
9806	1975	2.836275e+07	1.907006e+09	1.878643e+09	very_high
8889	1982	2.372625e+07	1.791694e+09	1.767968e+09	very_high
3	2015	1.839999e+08	1.902723e+09	1.718723e+09	very_high
8094	1995	3.148127e+07	1.583050e+09	1.551568e+09	very_high
10110	1961	2.917944e+07	1.574815e+09	1.545635e+09	very_high
7309	1980	4.762866e+07	1.424626e+09	1.376998e+09	very_high

```
In [92]: # calculate the profit mean value in the category
top_10_profit['profit'].mean()
```

```
Out[92]: 1953865824.9082134
```

```
In [93]: # compare the difference between the average of top 10 movies's revenue and low_budge
top_10_profit['profit'].mean()-low_budget_very_high_revenue['profit'].mean()
```

```
Out[93]: 1434663130.8359647
```

Conclusion: >- Even though some low_budget movies made very_high revenues, the difference between the average revenue for top_10_profit movies and low_budget_very_high_revenue movies are huge. >- This can imply that for the movies made with huge profit, they are made with huge budget too. We can not expect a movie with low budget to make extremely high profit; however it is possible to make low budget movies with moderately high profit, but the chances are not significant, there are only 61 very_high revenue movies in the low_budget table.

1.3.5 Question 5. What are the top 10 rated movies? and how is their profitability?

Since some movies have more vote_count, we can not directly compare a movie rated 10 with only 3 counts to the movie rated 7 with 100 counts. We will use IMDB's definition to calculate weighted average for rating score.

```
In [94]: # m is the minimum votes required to be listed in the chart;
m= tmdb['vote_count'].quantile(0.9)
m
```

```
Out[94]: 1371.7000000000003
```

```
In [95]: # C is the mean vote across the whole report
C=tmdb['vote_count'].mean()
print(C)
```

```
527.7202906071614
```

```
In [96]: # Create a table for top 10% highest rated movies
q_movies = tmdb.copy().loc[tmdb['vote_count'] >= m]
q_movies.shape
```

```
Out[96]: (386, 13)
```

```
In [97]: def weighted_rating(x, m=m, C=C):
          v = x['vote_count']
          R = x['vote_average']
          # Calculation based on the IMDB formula
          return (v/(v+m) * R) + (m/(m+v) * C)

          # Define a new feature 'score' and calculate its value with `weighted_rating()`
          q_movies['score'] = q_movies.apply(weighted_rating, axis=1)
```

```
In [98]: # show the top 10 rated movies
          q_movies.sort_values('score', ascending=False).head(10)
```

```
Out[98]:
```

	popularity	original_title \
2912	1.499784	Cloverfield
2643	2.449323	The Mummy Returns
6980	2.175284	Ocean's Twelve
7884	2.484654	Ghostbusters
21	5.337064	Southpaw
658	3.813740	Exodus: Gods and Kings
3416	1.499109	Rango
6631	0.890909	The Pursuit of Happyness
6578	1.603140	Blood Diamond
10094	0.142486	Home Alone

	cast	director \
2912	Lizzy Caplan Jessica Lucas Odette Annable Mich...	Matt Reeves
2643	Brendan Fraser Rachel Weisz John Hannah Arnold...	Stephen Sommers
6980	George Clooney Brad Pitt Catherine Zeta-Jones ...	Steven Soderbergh
7884	Bill Murray Dan Aykroyd Sigourney Weaver Harol...	Ivan Reitman
21	Jake Gyllenhaal Rachel McAdams Forest Whitaker...	Antoine Fuqua
658	Christian Bale Joel Edgerton John Turturro Aar...	Ridley Scott
3416	Johnny Depp Isla Fisher Ned Beatty Bill Nighy ...	Gore Verbinski
6631	Will Smith Jaden Smith Thandie Newton Brian Ho...	Gabriele Muccino
6578	Leonardo DiCaprio Djimon Hounsou Jennifer Conn...	Edward Zwick
10094	Macaulay Culkin Joe Pesci Daniel Stern John He...	Chris Columbus

	runtime	genres	vote_count \
2912	85	Action Thriller Science Fiction	1373
2643	130	Action Adventure Drama Fantasy Horror	1372
6980	125	Thriller Crime	1376
7884	107	Fantasy Action Comedy Science Fiction Family	1383
21	123	Action Drama	1386
658	153	Adventure Drama Action	1377
3416	107	Animation Comedy Family Western Adventure	1385
6631	117	Drama	1392
6578	143	Drama Thriller Action	1394

10094	103		Comedy Family	1393
-------	-----	--	---------------	------

	vote_average	release_year	budget_adj	revenue_adj	profit \
2912	6.4	2008	2.531967e+07	1.729475e+08	1.476279e+08
2643	5.8	2001	1.206858e+08	5.332507e+08	4.125649e+08
6980	6.4	2004	1.269890e+08	4.187685e+08	2.917795e+08
7884	7.2	1984	6.297126e+07	6.196634e+08	5.566921e+08
21	7.3	2015	2.759999e+07	8.437300e+07	5.677302e+07
658	5.6	2014	1.289527e+08	2.468817e+08	1.179290e+08
3416	6.5	2011	1.308687e+08	2.382049e+08	1.073362e+08
6631	7.5	2006	5.949180e+07	3.321560e+08	2.726642e+08
6578	7.2	2006	1.081669e+08	1.848334e+08	7.666646e+07
10094	7.0	1990	3.004017e+07	7.955384e+08	7.654982e+08

	revenue_level	score
2912	high	266.936686
2643	very_high	266.731612
6980	very_high	266.652226
7884	very_high	266.392537
21	low	266.160831
658	high	266.156773
3416	high	265.852803
6631	very_high	265.699578
6578	high	265.361653
10094	very_high	265.354260

- Above are the movies with highest rating score, as we can see, most have high or very_high revenue too, but it is different from the top_10_profit movies.

```
In [99]: q_movies['revenue_level'].value_counts()
```

```
Out[99]: very_high    242
         high         75
         low          64
         very_low      5
         Name: revenue_level, dtype: int64
```

```
In [100]: # Count for percentage of the one with low or very_low revenue
          (64+5)/q_movies.shape[0]
```

```
Out[100]: 0.17875647668393782
```

- About 17% of the highly rated movies have low revenue, most have very_high revenue.
- This can confirm that a good rating score can yield high revenue.

Conclusions

In this project we did comprehensive analysis on the movie database with a focus on movie revenues and other properties like genres and rating score. We did initial data exploration and answered all the questions

Summarize some of the featured findings: >- In general, higher budget can yield higher revenues; movies made with low budget can have moderately high revenue, but successful rate is not very high. >- Popular movies have more vote_count and also have higher revenues, and newer movies are more popular >- Although much more movies are produced over time, movies industry is getting more stable and annual average profit is lower compared with movies made 3 decades ago. >- 'Action', 'Adventure', 'Comedy', 'Drama', 'Thriller' are the most common genres for movies with very_high revenue >- Movies with high rating scores also have high revenues, but not necessarily yield the highest or extremely high profit.

Limitation of the project >- The dataset we are using contains some incorrect information. As we see in data cleaning process, more than half of the data is deleted since it contains 0 value for runtime, budget or revenue. If we can have a more complete data, the analysis can be more accurate. >- Since the project is main focusing on movie revenue analysis, and we find popularity, budget and rating score can have impact on revenue, but these information is not enough to make revenue prediction as there might be other factors that can affect movie revenue but not included in this dataset.